



# ktap introduction

Edited by:

ktap project

[www.ktap.org](http://www.ktap.org)

September 16, 2013

# Contents

- Introduction
- Architecture
- Building and Running
- Examples
- Performance
- More information

# ktap

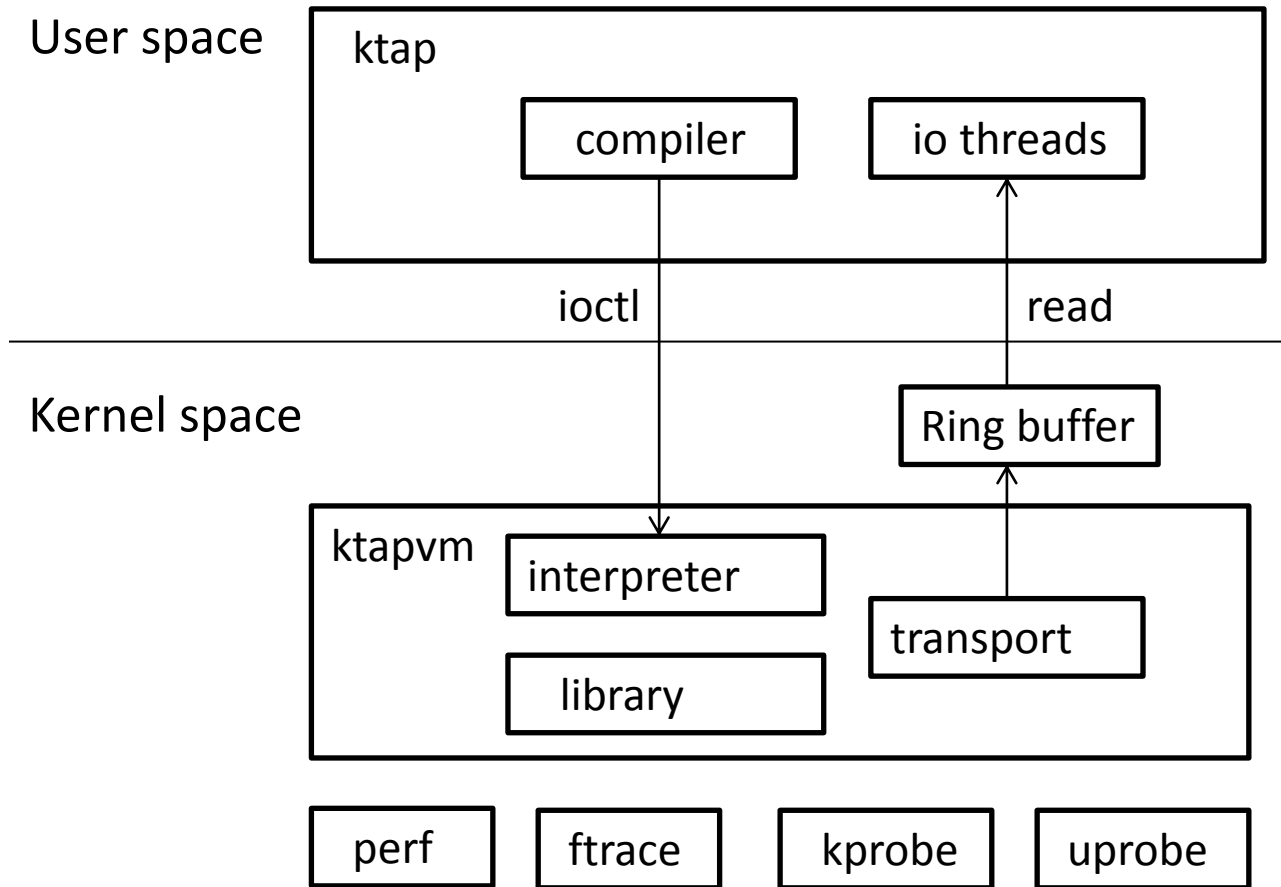
ktap is a new scripting dynamic tracing tool for Linux, it is designed to give operational insights that allow users to tune, troubleshoot and extend kernel and application.

ktap is similar with Systemtap and Dtrace, but it's not a clone of them.

# Feature highlights

- simple but powerful scripting language
- register based interpreter (heavily optimized) in kernel
- small and lightweight (6KLOC of interpreter)
- not depend on gcc for each script running
- easy to use in embedded system without debug info
- support for tracepoint, kprobe, uprobe, ftrace, timer, etc
- supported in x86, arm, ppc, mips
- good extendibility
- safety in sandbox

# Architecture



# Building and Running

1) Clone ktap from github

# git clone <http://github.com/ktap/ktap.git>

2) Compiling ktap

# cd ktap

# make      #generate ktapvm kernel module and ktap binary

3) Load ktapvm kernel module(make sure debugfs mounted)

# make load   #need to be root or have sudo access

4) Running ktap

# ./ktap scripts/helloworld.kp

# Examples

## 1) simplest one-liner command to enable all tracepoints

```
ktap -e "trace *:* { print(argevent) }"
```

## 2) syscall tracing on target process

```
ktap -e "trace syscalls:* { print(argevent) }" - ls
```

## 3) function tracing

```
ktap -e "trace ftrace:function { print(argevent) }"
```

```
ktap -e "trace ftrace:function /ip == mutex*/ { print(argevent) }"
```



# Examples (Cont.)

## 4) simple syscall tracing

```
#scripts/syscalls/syscalls.kp
trace syscalls:* {
    print(cpu(), pid(), execname(), argevent)
}
```

## 5) syscall tracing in histogram style

```
#scripts/syscalls/syscalls_count.kp
hist = {}

trace syscalls:sys_enter_* {
    hist[argname] += 1
}

trace_end {
    histogram(hist)
}
```

# Examples (Cont.)

## 6) kprobe tracing

```
#scripts/io/kprobes-do-sys-open.kp
trace probe:do_sys_open dfd=%di fname=%dx flags=%cx mode=+4($stack) {
    print("entry:", execname(), argevent)
}

trace probe:do_sys_open%return fd=$retval {
    print("exit:", execname(), argevent)
}
```

## 7) timer

```
tick-1ms {
    printf("time fired on one cpu\n");
}

profile-2s {
    printf("time fired on every cpu\n");
}
```

# Example: eventcount.kp

```
[root@localhost ktap]# ./ktap scripts/tracepoints/eventcount.kp
```

```
Press Control-C to stop.
```

```
^C
```

value	Distribution	count
rcu_utilization  @@@@@@@@		31106
cpu_idle  @@@@@@@		29738
hrtimer_start  @@		10087
hrtimer_cancel  @@		9881
softirq_entry  @@		9742
softirq_exit  @@		9722
softirq_raise  @@		9609
hrtimer_expire_entry  @		7823
hrtimer_expire_exit  @		7813
local_timer_entry  @		7121
local_timer_exit  @		7109
sched_switch  @		6653
kfree  @		4415
sched_wakeup		3576
sched_stat_runtime		3246
sched_stat_wait		2623
sched_stat_sleep		2018
tick_stop		1410
timer_start		418
timer_cancel		410
...		

# Example: schedtimes.kp

```
[root@localhost ktap]# ./ktap scripts/schedule/schedtimes.kp
```

```
^C
```

execname:	pid	run(us)	sleep(us)	io_wait(us)	queued(us)	total(us)
swapper/10:	0	2096967	0	0	3205	2100172
rcuos/10:	32	28	2096955	0	24	2097007
rcuos/11:	33	102	2078765	0	18126	2096993
kworker/0:2:	127	494	1980554	0	77	1981125
kworker/0:1H:	254	23	256115	0	23	256161
irqbalance:	534	1188	2032680	0	83	2033951
ksoftirqd/2:	39	21	981378	0	125	981524
kworker/11:2:	416	39	2000782	0	183	2001004
sshd:	8819	548	24115	0	144	24807
rcu_sched:	21	460	2078197	0	18428	2097085
ktap:	9074	24398	2074282	0	1709	2100389
kworker/10:0:	8734	175	1926607	0	200	1926982
rcuos/2:	24	266	2032665	0	660	2033591
ksoftirqd/11:	75	23	24076	0	62	24161
scsi_eh_0:	121	461	259339	0	61	259861
kworker/1:1:	89	279	260005	966	143	260427
kworker/2:1:	90	636	980666	0	270	981572
ktap:	9075	683	2096408	0	463	2097554
kworker/1:0:	8795	127	259986	0	14	260127
rcuos/0:	<u>22</u>	26	2096901	0	20	2096947

# Example: function\_profiler.kp

```
[root@localhost ktap]# ./ktap scripts/profiling/function_profiler.kp
```

```
^C
```

value	----- Distribution -----	count
read_hpet  @		74122
_raw_spin_lock  @		68179
ktime_get  @		58394
idle_cpu  @		49016
lookup_address  @		40932
get_vtime_delta  @		39740
notifier_call_chain		28927
source_load		24641
_raw_spin_unlock_irqrestore		22345
_raw_spin_lock_irqsave		22296
rcu_eqs_exit_common.isra.29		21572
rcu_eqs_enter_common.isra.28		21559
atomic_notifier_call_chain		21483
acct_account_cputime		20817
__acct_update_integrals		20817
cpuacct_account_field		20817
__vtime_account_system		20741
account_system_time		20741
vtime_account_irq_enter		19395
vtime_account_irq_exit		19382
...		

# Performance: boot time

The execution time of helloworld script is nearly **5x** than Systemtap

```
[root@localhost ktap]# time /usr/bin/stap -e 'probe begin { println("hello world") exit()}'  
hello world
```

real	0m0.575s
user	0m0.224s
sys	0m0.068s

```
[root@localhost ktap]# time ./ktap -e 'print("hello world") exit()'  
hello world
```

real	0m0.117s
user	0m0.001s
sys	0m0.009s

In Systemtap, normally you will need to wait **10+ seconds** before start to run script, for some complicated scripts, 20+ seconds maybe need to wait.

In ktap, you just need to wait **10+ms**(6ms to compile, 4ms to boot) to starting most scripts.

# Performance: function\_profiler

Idle:

```
%Cpu(s):  0.0 us,  0.2 sy,  0.0 ni, 99.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem:  2050668 total,  316800 used,  1733868 free,    27324 buffers
KiB Swap:  2113532 total,      0 used,  2113532 free,   121528 cached
```

Enable all kernel function profiling in ktap:

```
%Cpu(s):  0.0 us,  0.8 sy,  0.0 ni, 99.0 id,  0.0 wa,  0.2 hi,  0.0 si,  0.0 st
KiB Mem:  2050668 total,  334540 used,  1716128 free,    27324 buffers
KiB Swap:  2113532 total,      0 used,  2113532 free,   121528 cached
```

Enable all kernel function profiling in Systemtap:

```
%Cpu(s):  0.0 us,  9.0 sy,  0.0 ni, 91.0 id,  0.0 wa,  0.1 hi,  0.0 si,  0.0 st
KiB Mem:  2050668 total,  772996 used,  1277672 free,    54996 buffers
KiB Swap:  2113532 total,      0 used,  2113532 free,   273584 cached
```

kernel crashed after 5 seconds:

```
localhost login: [26958.622191] BUG: unable to handle kernel paging request at ffff88007c3ceff8
[26958.622191] IP: [<ffffffff81035960>] ftrace_update_ftrace_func+0xa0/0xa0
[26958.622191] PGD 1fb4067 PUD 1fb7067 PMD 36e15063 PTE 800000007c3ce161
[26958.622191] Oops: 0003 [#1] SMP
[26958.622191] Modules linked in: stap_34352c5bf185b5341ceb4370930d0024_4_10083(OFF) ktapvm(OFF) ebttable_1
mangle ip6t_REJECT nf_conntrack_ipv6 nf_defrag_ipv6 iptable_nat nf_nat_ipv4 nf_nat iptable_mangle nf_co
bles joydev virtio_balloon mperf e1000 [last unloaded: stap_e68c0f0ec1a89e8f3b5f9c59cea24495_9867]
[26958.622191] CPU: 0 PID: 10083 Comm: stapio Tainted: GF      W O 3.11.0-rc4+ #2
[26958.622191] Hardware name: Bochs Bochs, BIOS Bochs 01/01/2007
[26958.622191] task: ffff8800762eaf20 ti: ffff8800730a2000 task.ti: ffff8800730a2000
[26958.622191] RIP: 0010:[<ffffffff81035960>] [<ffffffff81035960>] ftrace_update_ftrace_func+0xa0/0xa0
```

# Performance: stack\_profiler

System stack sample with 1ms period

Idle:

```
%Cpu(s):  0.0 us,  0.1 sy,  0.0 ni, 99.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem:  2050668 total,  289936 used,  1760732 free,    344 buffers
KiB Swap:  2113532 total,    0 used,  2113532 free,   31188 cached
```

Enable stack profiling in ktop:

```
%Cpu(s):  0.0 us,  0.5 sy,  0.0 ni, 99.4 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem:  2050668 total,  307172 used,  1743496 free,    408 buffers
KiB Swap:  2113532 total,    0 used,  2113532 free,   31612 cached
```

Enable stack profiling in Systemtap:

```
%Cpu(s):  0.0 us,  1.2 sy,  0.0 ni, 98.8 id,  0.0 wa,  0.0 hi,  0.0 si,  0.0 st
KiB Mem:  2050668 total,  399756 used,  1650912 free,   18524 buffers
KiB Swap:  2113532 total,    0 used,  2113532 free,   55468 cached
```



# More information

- Project status:
  - Released 0.2
  - only support kernel newer than 3.1
- Project home: [www.ktap.org](http://www.ktap.org)
- Code location: <https://github.com/ktap/ktap.git>
- Mailing list: [ktap@freelists.org](mailto:ktap@freelists.org)
- Documentation: `ktap/doc/`
- Sample scripts: `ktap/scripts/`